



THE DATA & AI CHALLENGE · REDROB × INDIA RUNS

Parakh

/pə-rəkh/ — to assay; to test whether gold is genuine.

A candidate ranker that **verifies** each profile before it trusts it — telling genuine engineers from convincing fakes.

Team Evolve

Rayyan Ahmed Shaikh · Inchara · Harshita Nagesh

THE PROBLEM

Great candidates are missed — not absent, just unseen.

Recruiters scan hundreds of profiles and still miss the right person — because **keyword filters can't see fit**.

They reward whoever lists the most buzzwords, and overlook the engineer who **built the exact system** but described it plainly.

In a pool of 100,000 the real fits are a needle in a haystack — and half the score rides on the top ten.

100,000

profiles in the pool

~ few hundred

genuinely fit the role

Top 10

carry 50% of the score

The dataset is engineered to punish keyword search

● Keyword stuffers

Lists every AI skill; the actual job is Marketing or Ops. Counting keywords ranks the fake at the top.

● Plain-language gems

7,334 people describe real retrieval / ranking work under a non-ML title. Title filters miss them entirely.

● Behavioural twins

Near-identical on paper; they differ only in engagement signals — who actually responds and shows up.

● Honeypots (~80)

Impossible profiles (expert skills, 0 months' use). Rank >10% of them in the top 100 and you're disqualified.

What most teams will submit — and why it loses

The obvious pipeline

1. Embed the job description
2. Embed all 100,000 profiles
3. Cosine similarity → sort

What the traps do to it

- Keyword-stuffers rank #1 — buzzwords beat real work
- Plain-language gems get buried below them
- Impossible honeypots slip into the top 10
- → **honeypot rate breaks 10% → disqualified**

The dataset was built to beat exactly this. So we don't play that game.

Assay, don't match.

Like an assayer testing gold, Parakh **verifies every profile before it trusts a word of it** — then ranks only on evidence that survives.

-  **Read the work, not the words**
Evidence in the career story outweighs any skills list.
-  **Reject the impossible**
Fakes and honeypots are removed before ranking, not after.
-  **Trust what judges agree on**
The very top is where independent AI graders converge.
-  **Prove it, don't hope**
Every design choice is measured against an independent judge.

An LLM's judgment — at CPU speed

The hard rule: the ranking step must run offline, CPU-only, no LLM calls, under 5 minutes over 100,000 profiles. An LLM call per candidate simply cannot fit — or scale in production.

Think — offline

Two frontier LLMs (DeepSeek-V4, Qwen-235B) grade every plausible candidate on the job rubric. Slow, brilliant, unconstrained.



Bake — into artifacts

Its verdicts are frozen into a small cached file that ships with the code — a precomputed feature.



Serve — in milliseconds

The offline ranker reads those verdicts. No LLM, no GPU, no network at run-time.

The constraint becomes the advantage: frontier intelligence, distilled into a ranker that's fast, legal, and reproducible.

Two phases — all the power is offline

Pre-compute

offline · LLM + GPU allowed

Sweep all 100k
+ features



DeepSeek-V4 + Qwen-235B
grade fit 0–5



GLM-5.2
3rd judge

Ranking step

offline · CPU · no network · < 5 min

Integrity Gate



Score all
100,000



Top 100
+ reasons

Spec §3 compliance: no LLM, GPU or network during ranking — every model call lives in pre-compute and ships as a cached artifact the ranker reads.

Reject the impossible and the fake — first



Honeypot check

'Expert' in a skill used 0 months · tenure longer than the whole career · contradictory dates. ~84 hits ≈ the ~80 planted honeypots.



Stuffer check

Many AI skills, but the career narrative shows none of the work, and the title is wrong. Keywords without evidence don't count.

VERDICT AT RUNTIME

Flagged profiles are forced to the bottom
— they can never reach the top 100.

0

honeypots in our top 100



WHY IT WORKS — A CONCRETE CASE

Seeing through the disguise

THE DISGUISE

Marketing Manager · 9 yrs

Skills: RAG, NLP, Embeddings, Ranking

Career: brand campaigns & strategy. No ML system ever built.

Judge verdict: TIER 1 — keywords, no evidence

THE REAL THING

Data Engineer · 7 yrs

No 'RAG' or 'Pinecone' in the skills list.

Career: “built and shipped a product recommendation system at scale.”

Judge verdict: TIER 5 — real work, plainly told

Keyword counting ranks the fake #1 and never finds the gem. Parakh reads the work and flips both.



What the teacher actually judges

Must-haves it rewards

- Production retrieval / ranking / recsys, shipped to users
- Embeddings + a vector DB (FAISS, Pinecone, Qdrant...)
- Rigorous eval — NDCG, MRR, A/B testing
- 5–9 yrs at product companies; in / open to India

Disqualifiers it catches

- Stuffers — AI keywords with no supporting work
- Pure research / academia, no production
- Recent LangChain-only 'AI experience'
- Whole career at IT-services / consulting

Grades every plausible candidate 0–5 with a written reason — the rubric a senior recruiter carries in their head.

Three model families — the top is where they agree

DeepSeek + Qwen

two independent teachers

GLM-5.2

3rd judge — contested top



Average the tiers

Only candidates **every model family** endorses rise to the very top.

Why it pays off

Different model families make different mistakes. Requiring agreement cancels each one's overconfidence.

An echo of my *Quorum* project: trust rises where independent models converge.



Beyond the résumé

Availability signals

- Recruiter response rate — will they reply?
- Last active — a 6-month ghost isn't hireable
- Open-to-work + notice period
- Interview completion, saved-by-recruiters

Red flags it down-weights

- Based abroad and won't relocate
- Long notice period
- Title-chaser — job-hops every ~18 months
- 'Architect' who hasn't coded in 18 months

A perfect-on-paper candidate who never replies isn't a hire. These signals sharpen the very top.



RIGOR

We didn't fly blind

No public leaderboard and only three submissions — so most teams guess. We built our own eval harness and judged every version with a **third, independent model** that never touches the ranking.



DeepSeek + Qwen

two teachers — 5,500 graded each



GLM-5.2

3rd judge — contested top




Llama-3.3-70B

independent scorer

Rule: a change shipped only if the independent judge scored it higher. Nothing made the cut on faith.

Every layer earned its place

NDCG@10, scored by an independent third model (Llama-3.3-70B)



10 / 10

top-10 rated tier-5 by three model families

0.98

NDCG@50 (ensemble)

Honest read: *three different models agree the top-10 is excellent — a strong proxy, not the organisers' hidden ground truth.*

Why this is different

Everyone else

- ✗ Embed everything, cosine, sort
- ✗ Fooled by keyword-stuffers
- ✗ Buries plain-language gems
- ✗ Risks honeypot disqualification
- ✗ Flies blind — no way to measure

Parakh / Team Evolve

- ✓ Assays every profile before trusting it
- ✓ Distills a frontier LLM into a legal offline ranker
- ✓ Two independent judges must agree at the top
- ✓ 0 honeypots · top-10 all tier-5
- ✓ Proven by a third, independent model

Same rules as everyone. A fundamentally different bet: understand the candidate, don't pattern-match the keywords.



Parakh

Assaying talent — genuine fits, surfaced and trusted.

```
python rank.py --candidates candidates.jsonl
```

Team Evolve

Rayyan Ahmed Shaikh · Inchara · Harshita Nagesh